

# Index

## A

### A/B testing

canary releases, [Considerations When Sending Models to Production](#)

considerations in MLOps context, [A/B testing](#)

of new and existing model versions, [The Feedback Loop](#)

performance monitoring for marketing recommendation engine, [Monitoring Performance](#)

for recommendation engines using collaborative filtering, [The Role of Machine Learning](#)

using in online evaluation of models, [Online Evaluation](#), [A/B testing](#)

### accountability

GxP guidelines focus on, [Pharmaceutical Regulation in the US: GxP in Responsible AI](#), [MLOps for Responsible AI](#), [The Emergence of Responsible AI](#)

### accuracy, precision, and recall, [Prepare for Production](#)

metrics collected in preproduction model testing, [Key Testing Considerations](#)

### adversarial attacks, [Adversarial Attacks](#)

### AI, [The Emergence of Responsible AI](#)

(see also Responsible AI)

new wave of AI-specific regulations, [The New Wave of AI-Specific Regulations-The Emergence of Responsible AI](#)

Responsible AI, MLOps for, [MLOps for Responsible AI](#)

### AIOps versus MLOps, [Why Now and Challenges](#)

### algorithms (machine learning), [A Primer on Machine Learning](#), [Required Components](#)

computing power considerations, [Different ML Algorithms](#), [Different MLOps Challenges](#)

MLOps considerations by algorithm type, [Different ML Algorithms](#), [Different MLOps Challenges](#)

online learning, [How Often Should Models Be Retrained?](#)

requirement for tabular input data, [Feature Engineering Techniques](#)

smoothness of predicted curve, [Modeling](#)

analytics use cases, understanding and classifying, [Step 1: Understand and Classify the Analytics Use Cases](#)

anomaly detection, [Model Misbehavior](#)

anonymizing or pseudo-anonymizing data, [Data Governance](#)

Apache Spark, [Categories of Model Deployment](#)

APIs

marketing recommendation engine API endpoints, [Pipeline Structure and Deployment Strategy](#)

REST API for model-as-a-service or live-scoring model, [Model Deployment Types and Contents](#)

ARIMA (autoregressive integrated moving average), [Modeling, Monitoring](#)

artifacts (ML), [Building ML Artifacts-The Testing Pipeline](#)

assumptions (model), [Version Management and Reproducibility](#)

auditability, [The New Wave of AI-Specific Regulations, Element 4: Model Management at Scale](#)

aiding with QA for machine learning, [Reproducibility and Auditability](#)

and reproducibility, [Reproducibility and Auditability](#)

automation

automated feature selection, [How Feature Selection Impacts MLOps Strategy](#)

automated model deployment, [Model Deployment Requirements](#)

automated model documentation, [Responsible AI](#)

automated model packaging and delivery, [People of MLOps, Data Scientists](#)

automated reporting tools on all models, [People of MLOps](#)

automatically controlling models, marketing recommendation system, [Option to Automatically Control Models](#)

in experimentation during model development, [Experimentation](#)

feedback loop, [Deploy to Production](#)

of tests in testing pipeline, [The Testing Pipeline](#)

of versioning and reproducibility tasks, [Version Management and Reproducibility](#)

## B

batch scoring, [Categories of Model Deployment](#)

volume of data becoming too large, distributing computation,

[Scaling Deployments](#)

Bayesian tests, [The Feedback Loop](#)

biases

analyzing models for fairness, [Key Testing Considerations](#)

inappropriate biases in models, [Data Governance](#)

introduced by ground truth monitoring, [Ground Truth Evaluation](#)

in ML black box judgments, [Who Decides What Governance the Organization Needs?](#)

model bias considerations in consumer credit risk management,

[Model Bias Considerations](#)

reputational risk due to, [The Origins of ML Model Risk](#)

Responsible AI position on, [Element 2: Bias](#)

sample selection bias, [Example Causes of Data Drift](#)

tuning bias/variance trade-off, [Experimentation](#)

black box models, [Impact of Responsible AI on Modeling](#)

blue-green deployment, [Model Deployment Types and Contents](#),  
[Considerations When Sending Models to Production](#)

bounds for retraining frequency, [How Often Should Models Be Retrained?](#)

business concerns in monitoring ML models, [Business Concerns](#)

business decision modeling, [Subject Matter Experts](#)

business metrics for model performance monitoring, [Ground Truth Evaluation](#), [Univariate statistical tests](#)

business objectives, establishing for ML model, [Establishing Business Objectives](#)

business use case, consumer credit risk management application,  
[Background: The Business Use Case](#)

## C

canary releases, [Considerations When Sending Models to Production](#)

CCPA (California Consumer Privacy Act), [Governance](#), [GDPR and CCPA Data Privacy Regulations](#)

champion/challenger, [Key Testing Considerations](#),  
[Champion/Challenger](#)

Chi-squared test, [Univariate statistical tests](#)

CI/CD pipelines, [CI/CD Pipelines-CI/CD Pipelines](#)

continuous delivery for end-to-end machine learning process, [The Feedback Loop](#)

DevOps role in managing, [DevOps](#)

for different models, [Scaling Deployments](#)

Jenkins build system, [The Testing Pipeline](#)

ML artifacts in, [What's in an ML Artifact?](#)

robust, for model deployments, [Model Deployment Requirements](#)

collaborative filtering, [The Role of Machine Learning](#)

compression techniques, use in model definition optimization,  
[Performance considerations](#)

computational metrics from model testing, [Key Testing Considerations](#)

computing power

ML model development and, [Different ML Algorithms](#), [Different MLOps Challenges](#)

required for inference on ML models, [Final Thoughts on Runtime Environments](#)

concept drift, [Input Drift Detection](#)

conformal prediction, [Model Misbehavior](#)

consumer credit risk management, [MLOps in Practice: Consumer Credit Risk Management-Closing Thoughts](#)

business use case, [Background: The Business Use Case](#)

problem definition and data acquisition, [Background: The Business Use Case](#)

deploying model to production, [Deploy to Production](#)

model development, [Model Development](#)

bias considerations, [Model Bias Considerations](#)

preparing model for production, [Prepare for Production](#)

consumption forecast for power grid, [MLOps in Practice: Consumption Forecast-Closing Thoughts](#)

data collection, [Data Collection-Data Collection](#)

deployment of models, [Deployment](#)

implementation, [Implementation-Modeling](#)

modeling, [Modeling](#)

monitoring, [Monitoring](#)

power systems, [Power Systems-Power Systems](#)

problem definition, using machine learning or not, [Problem Definition: Machine Learning, or Not Machine Learning?-Spatial and Temporal Resolution](#)

spatial and temporal resolution, [Spatial and Temporal Resolution](#)

containerization, [Containerization-Containerization](#)

solving problems of dependencies in ML model deployments, [Model Deployment Types and Contents](#)

continuous integration and continuous delivery (see CI/CD pipelines)  
costs

for algorithms that train themselves, [How Often Should Models Be Retrained?](#)

for retraining models versus performance improvement, [How Often Should Models Be Retrained?](#)

cross-trained models, [How Often Should Models Be Retrained?](#)

curse of dimensionality, [Model Misbehavior](#)

customizability, marketing recommendation engine model, [Scalability and Customizability](#)

## D

dark launch, [Champion/Challenger](#)

(see also champion/challenger)

data

collection for consumption forecast for power grid, [Data Collection-Data Collection](#)

considerations in Responsible AI, [Element 1: Data](#)

customer data, preparation for marketing recommendation engine, [Data Preparation](#)

processing for consumption forecast system, [Implementation](#)

reproducibility, [Version Management and Reproducibility](#)

data access before validation and launch to production, [Data Access Before Validation and Launch to Production](#)

data architects, [Machine Learning Architect](#)

data cleansing, [Feature Engineering and Selection](#)

data drift, [Input Drift Detection](#)

for consumer credit risk management model, [Prepare for Production](#)

example causes of, [Example Causes of Data Drift](#)

data engineers, [Data Engineers](#)

data exploration, [Data Exploration](#)

data governance, [Data Governance](#)

questions for ML model data sources, [Data Sources and Exploratory Data Analysis](#)

data pipeline structure for recommendation engine project, [Pipeline Structure and Deployment Strategy](#)

data privacy, [Governance](#), [Current Regulations Driving MLOps Governance](#)

GDPR and CCPA regulations on, [Governance](#), [GDPR and CCPA Data Privacy Regulations](#)

data scientists, [Data Scientists-Data Scientists](#)

collaboration with SMEs in ML model life cycle, [Subject Matter Experts](#)

concerns in ML model monitoring, [Data Scientist Concerns](#)

ground truth, [Ground truth](#)

input drift, [Input drift](#)

role in and needs from MLOps, [Data Scientists](#)

role in machine learning model life cycle, [Data Scientists](#)

data sources for machine learning models, [Data Sources and Exploratory Data Analysis](#)

DataOps, [Defining MLOps and Its Challenges](#)

decision modeling (business), [Subject Matter Experts](#)

decision-making processes, statistically driven, [Who Decides What Governance the Organization Needs?](#)

deep learning, [Different ML Algorithms](#), [Different MLOps Challenges](#), [Feature Engineering Techniques](#), [Impact of Responsible AI on Modeling](#)

degradation of model performance

common approaches to discovering, [Data Scientist Concerns](#)

understanding, [Understanding Model Degradation-Input Drift Detection](#)

delivery, [Deployment Strategies](#)

(see also CI/CD pipelines)

continuous delivery versus deployment, [Deployment Strategies](#)  
dependencies

partial dependency plots, [Responsible AI](#)

on production environment, reducing, [Model Deployment Types and Contents](#)

deployment strategies, [Deployment Strategies-Maintenance in Production](#)

categories of model deployment, [Categories of Model Deployment](#)

concepts and terminology, [Deployment Strategies](#)

considerations in sending models to production, [Considerations When Sending Models to Production](#)

maintenance of models in production, [Maintenance in Production](#)

deployments

broader model deployment, greater risks from, [Model Risk Mitigation](#)

consumption forecast models, [Deployment](#)

deploying to production, [Deploying to Production-Closing Thoughts](#)

building ML artifacts, [Building ML Artifacts-The Testing Pipeline](#)

CI/CD pipelines, [CI/CD Pipelines](#)

consumer credit risk management model, [Deploy to Production](#)

containerization, [Containerization-Containerization](#)

scaling deployments, [Scaling Deployments-Requirements and Challenges](#)

strategies for, [Deployment Strategies-Maintenance in Production](#)

deployment, defined, [Deployment Strategies](#)

marketing recommendation engine model, [Model Training and Deployment-Ability to Turn Recommendations On and Off](#)

model deployment types and contents, [Model Deployment Types and Contents](#)

model deployment requirements, [Model Deployment Requirements](#)

development environments, adaptation to production environments, [Adaptation from Development to Production Environments-Performance considerations](#)

DevOps, [DevOps](#)

concerns in ML model monitoring, [DevOps Concerns](#)



MLOps and, [Defining MLOps and Its Challenges](#)  
monitoring of ML models, [Monitoring and Feedback Loop](#)  
role in and needs from MLOps, [DevOps](#)  
role in machine learning model life cycle, [DevOps](#)  
DI (Data Integrity), [Pharmaceutical Regulation in the US: GxP](#)  
dimensionality, curse of, [Model Misbehavior](#)  
dimensioning constraints on model development, [Impact of Responsible AI on Modeling](#)  
disaggregations of data, [Implementation](#), [Monitoring](#)  
distances between probability distributions, [Prepare for Production](#)  
distillation (model), [Performance considerations](#)  
distributed computation, [Scaling Deployments](#)  
distribution of data, [Input Drift Detection](#)  
divergence between training and testing phases, [Drift Detection in Practice](#)  
Docker, [Containerization](#)  
deployment of models through, [Deploy to Production](#)  
using Kubernetes with, [Containerization](#)  
documentation of model development, [Responsible AI](#)  
domain knowledge, importance in data exploration for models, [Data Exploration](#)  
domains  
domain classifier, [Domain classifier](#)  
model retraining frequency and, [How Often Should Models Be Retrained?](#)  
drift, [Input Drift Detection](#)  
(see also input drift)  
detection in practice, [Drift Detection in Practice-Interpretation of results](#)  
example causes of data drift, [Example Causes of Data Drift](#)  
input drift detection techniques, [Input Drift Detection Techniques](#)  
measuring for consumer credit risk assessment model, [Prepare for Production](#)  
monitoring, [Monitoring](#)  
monitoring and mitigation measures, [Closing Thoughts](#)



## E

EDA (exploratory data analysis), [Data Sources and Exploratory Data Analysis](#), [Data Sources and Exploratory Data Analysis](#), [Data Exploration](#)

efficiency, data scientists' need for from MLOps, [Data Scientists](#) elastic systems, [Scaling Deployments](#)

embedded models, [Model Deployment Types and Contents](#)

embeddings, [Feature Engineering Techniques](#)

engaging and educating groups responsible for governance, [Step 7: Engage and Educate](#)

environmental, social, and governance (ESG) performance indicators, [Governance](#)

environments

changing rapidly, multiplying model risks, [Changing Environments](#)

information needed by data scientists, [Cross-Checking Model Behavior](#)

providing exact description for model version management,

[Containerization](#)

reproducibility and, [Version Management and Reproducibility](#)

ethical position, establishing in MLOps, [Step 2: Establish an Ethical Position](#)

EU

General Data Protection Regulation (GDPR), [Governance](#), [GDPR and CCPA Data Privacy Regulations](#)

key requirements for trustworthy AI applications, [The New Wave of AI-Specific Regulations](#)

evaluation datasets, [Required Components](#)

experimentation in model development, [Experimentation-Experimentation](#)

impacts on MLOps strategy, [Experimentation](#)

marketing recommendation engines, [Design and Manage Experiments](#)

explainability, [Responsible AI](#)

for decisions made by ML models affecting humans, [Impact of Responsible AI on Modeling](#)

in Responsible AI, [The Emergence of Responsible AI](#)

exploratory data analysis (see EDA)

## F

Facebook-Cambridge Analytica affair, [Who Decides What Governance the Organization Needs?](#)

fairness

reassuring public that ML is fair, [Who Decides What Governance the Organization Needs?](#)

requirements having dimensioning constraints on model development, [Impact of Responsible AI on Modeling](#)

subpopulation analysis and model fairness, [Key Testing Considerations](#)

feature drift, [Input Drift Detection](#)

feature stores, [How Feature Selection Impacts MLOps Strategy](#)

features, [Data Sources and Exploratory Data Analysis](#)

controlling feature-value intervals, [Model Misbehavior](#)

drift attributed to, use in mitigating impact of drift, [Interpretation of results](#)

engineering and selection, [Feature Engineering and Selection](#), [Feature Engineering and Selection-How Feature Selection Impacts MLOps Strategy](#)

feature engineering techniques, [Feature Engineering Techniques](#)

impacts of feature selection on MLOps strategy, [How Feature Selection Impacts MLOps Strategy](#)

statistical test on data from source and target distribution, [Univariate statistical tests](#)

federated learning approach to model retraining, [The Feedback Loop](#)

feedback loop, [The Feedback Loop-Closing Thoughts](#)

automating or not, [Deploy to Production](#)

infrastructure, main components of, [The Feedback Loop](#)

logging, [Logging](#)

model evaluation, [Model Evaluation-Model evaluation store](#)

logical model, [Logical model](#)

model evaluation store, [Model evaluation store](#)

online evaluation of models in production, [Online Evaluation-A/B testing](#)

A/B testing, [A/B testing](#)

champion/challenger, [Champion/Challenger](#)

financial crisis of 2007-2008, [Financial Model Risk Management Regulation](#)

financial model risk management regulation, [Financial Model Risk Management Regulation](#)

financial risk management techniques, [Model Development](#)

Flask framework, [Model Deployment Requirements](#)

forecasting, [MLOps in Practice: Consumption Forecast](#)

(see also consumption forecast for power grid)

forecasting techniques by use case, [Problem Definition: Machine Learning, or Not Machine Learning?](#)

spatial and temporal resolution of consumption, [Spatial and Temporal Resolution](#)

## G

GAM (generalized additive models), [Modeling](#)

GDPR (General Data Protection Regulation), [Governance](#), [GDPR and CCPA Data Privacy Regulations](#)

generalization capacity (models), [In Theory](#)

Git, [Building ML Artifacts](#)

Google, smartphone keyboard software, GBoard, [The Feedback Loop](#)

governance, [Governance-Process Governance](#), [Model Governance-Closing Thoughts](#)

AI-specific regulations, new wave of, [The New Wave of AI-Specific Regulations-The New Wave of AI-Specific Regulations](#)

status of AI governance initiatives worldwide, [The New Wave of AI-Specific Regulations](#)

application to MLOps, [Governance](#)

data governance, [Data Governance](#)

process governance, [Process Governance](#)

critical role in machine learning security, [Other Vulnerabilities](#)

current regulations driving MLOps governance, [Current Regulations Driving MLOps Governance-GDPR and CCPA Data](#)

## Privacy Regulations

financial model risk management, [Financial Model Risk](#)

## Management Regulation

GDPR and CCPA data privacy regulations, [GDPR and CCPA Data](#)

## Privacy Regulations

pharmaceutical regulation in US, GxP, [Pharmaceutical](#)

## Regulation in the US: GxP

key elements of Responsible AI, [The Emergence of Responsible AI-](#)

## Element 5: Governance

bias, [Element 2: Bias](#)

data, [Element 1: Data](#)

governance techniques, [Element 5: Governance](#)

inclusiveness, [Element 3: Inclusiveness](#)

model management at scale, [Element 4: Model Management at Scale](#)

matching with risk level, [Matching Governance with Risk Level](#)

template for MLOps, [A Template for MLOps Governance-Step 8: Monitor and Refine](#)

determining governance policies, [Step 4: Determine Governance Policies](#)

engaging end educating, [Step 7: Engage and Educate](#)

establishing an ethical position, [Step 2: Establish an Ethical Position](#)

establishing responsibilities, [Step 3: Establish Responsibilities](#)

integrating governance policies into MLOps process, [Step 5: Integrate Policies into the MLOps Process](#)

monitoring and refining governance, [Step 8: Monitor and Refine](#)

selecting tools for centralized governance management, [Step 6: Select the Tools for Centralized Governance Management](#)

understanding/classifying analytics use cases, [Step 1: Understand and Classify the Analytics Use Cases](#)

who decides on organization needs, [Who Decides What Governance the Organization Needs?](#)

gradient boosting algorithms, [Model Development](#)

ground truth, [Ground truth](#)

evaluation of, in monitoring of performance degradation, [Ground Truth Evaluation](#)

GxP pharmaceutical regulations in US, [Pharmaceutical Regulation in the US: GxP](#)

## H

health checks, [Maintenance in Production](#)  
health diagnosis or prediction example, [In Theory](#)  
high bias model (underfitting), [Experimentation](#)  
high variance model (overfitting), [Experimentation](#)  
historizing data, [Implementation](#)  
house price example, model prediction and generalization, [In Theory](#)  
human-centered approach, [The Emergence of Responsible AI](#)  
human-in-command, [Ability to Turn Recommendations On and Off](#)  
human-in-the-loop (HITL), [Element 3: Inclusiveness](#)  
hyperparameters of ML algorithms, [Required Components](#)

## I

impact coding, [Feature Engineering Techniques](#)  
implementation, reproducibility of models and, [Version Management and Reproducibility](#)  
inclusiveness, [Element 3: Inclusiveness](#)  
individual conditional expectation (ICE) computations, [Impact of Responsible AI on Modeling](#)  
inference, computing power required for, [Final Thoughts on Runtime Environments](#)  
input drift, [Input drift](#)  
    detection of, [Input Drift Detection](#)  
    detection techniques, [Input Drift Detection Techniques](#)  
        domain classifier, [Domain classifier](#)  
        interpretation of results, [Interpretation of results](#)  
        univariate statistical tests, [Input Drift Detection Techniques](#)  
    model retraining motivated by, [Iteration](#)  
integration, [Deployment Strategies](#)  
    (see also CI/CD pipelines)  
intentionality, Responsible AI, [MLOps for Responsible AI](#), [The Emergence of Responsible AI](#)

interactions between models, risks generated by, [Interactions Between Models](#)

interpretability (ML), bias detection through, [Element 2: Bias iterations of ML models, Iteration](#)

feedback loop, [The Feedback Loop](#)

iterating on model deployed to millions of devices, [The Feedback Loop](#)

## J

Java-based environment, model developed in Python, [Tooling considerations](#)

Jenkins build system, [The Testing Pipeline](#)

## K

Kolmogorov-Smirnov test, [Univariate statistical tests, Prepare for Production](#)

KPIs (key performance indicators), [Establishing Business Objectives, Matching Governance with Risk Level](#)

in business scoping, [Step 5: Integrate Policies into the MLOps Process](#)

for governance, [Step 8: Monitor and Refine](#)

measurement of, [Data Sources and Exploratory Data Analysis](#)

monitoring for deployed ML models, [Business Concerns](#)

subject matter experts contributing, [Subject Matter Experts](#)

tracking model performance against, [Step 6: Select the Tools for Centralized Governance Management](#)

training different models for different KPIs, [Option to Automatically Control Models](#)

Kubernetes, [Model Deployment Types and Contents, Containerization](#)

autoscaling properties, [Scaling Deployments](#)

using with Spark, [Scaling Deployments](#)

## L

labeled sample subset, problems in ground truth monitoring, [Ground Truth Evaluation](#)

latency

metrics on, collection in preproduction model testing, [Key Testing Considerations](#)

model conversions and, [Performance considerations](#)

linear algorithms, [Different ML Algorithms](#), [Different MLOps Challenges](#)

linear models, advanced, [Model Development](#)

Linux cgroups, [Model Deployment Requirements](#)

live-scoring model, [Model Deployment Types and Contents](#)

logging

activity logging of deployed model, [Deploy to Production](#)

component in ML feedback loop, [Logging](#)

using to generate centralized datasets for use by model designer, [Requirements and Challenges](#)

logical model, [Logical model](#)

comparing performance between different versions, [Model evaluation store](#)

versioning evolution of, [Model evaluation store](#)

LSTM (long short-term memory) neural networks, [Modeling](#)

## M

machine learning, [Developing Models](#)

(see also models)

continuous delivery for end-to-end process, [The Feedback Loop](#)

deciding whether to use or not, [Problem Definition: Machine Learning, or Not Machine Learning?-Spatial and Temporal Resolution](#)

primer, [A Primer on Machine Learning](#)

quality assurance for, [Quality Assurance for Machine Learning-Key Testing Considerations](#)

role in recommendation engines, [The Role of Machine Learning](#)

security issues, [Machine Learning Security-Other Vulnerabilities](#)

machine learning architects, [Machine Learning Architect](#)

role in and needs from MLOps, [Machine Learning Architect](#)

role in machine learning model life cycle, [Machine Learning Architect](#)



machine learning metrics, monitoring, [Maintenance in Production](#)

machine learning models (see models)

maintenance of models in production, [Maintenance in Production](#)

manual control of models, [Option to Manually Control Models](#)

MAPE (mean absolute percentage error), [Monitoring](#)

marketing data mining urban legend, [The Rise of Recommendation Engines](#)

marketing recommendation engines, [MLOps in Practice: Marketing Recommendation Engines-Closing Thoughts](#)

data preparation for, [Data Preparation](#)

designing and managing experiments on, [Design and Manage Experiments](#)

model training and deployment, [Model Training and Deployment-Ability to Turn Recommendations On and Off](#)

challenges in, [Model Training and Deployment](#)

monitoring and feedback, [Monitoring and Feedback-Monitoring Performance](#)

pipeline structure and deployment strategy, [Pipeline Structure and Deployment Strategy](#)

rise of recommendation engines, [The Rise of Recommendation Engines-Push or Pull?](#)

push or pull recommendations, [Push or Pull?](#)

role of machine learning, [The Role of Machine Learning](#)

MATLAB scripts, [Deployment](#)

minority populations, representation in ML model data sources, [Data Sources and Exploratory Data Analysis](#)

misbehavior of models, [Model Misbehavior](#)

missing value imputation, [Implementation](#)

MLOps

considerations with A/B testing, [A/B testing](#)

defining, and its challenges, [Defining MLOps and Its Challenges-Defining MLOps and Its Challenges](#)

governance template for, [A Template for MLOps Governance-Step 8: Monitor and Refine](#)

mitigating risks of ML models, [MLOps to Mitigate Risk-MLOps for Responsible AI](#)

versus ModelOps versus AIOps, [Why Now and Challenges](#)

real-world examples

consumer credit risk management, [MLOps in Practice: Consumer Credit Risk Management-Closing Thoughts](#)

consumption forecast, [MLOps in Practice: Consumption Forecast-Closing Thoughts](#)

marketing recommendation engines, [MLOps in Practice: Marketing Recommendation Engines-Closing Thoughts](#)

for scale, [MLOps for Scale](#)

model evaluation stores, [Model Evaluation](#), [Model evaluation store](#)

model risk manager/auditor, [Model Risk Manager/Auditor](#)

role in and needs from MLOps, [Model Risk Manager/Auditor](#)

role in machine learning model life cycle, [Model Risk Manager/Auditor](#)

model-as-a-service, or live-scoring model, [Model Deployment Types and Contents](#)

models

developing, [Model Development-Responsible AI](#), [Developing Models-Closing Thoughts](#)

consumer credit risk management model, [Model Development](#)

data exploration, [Data Exploration](#)

data sources and exploratory data analysis, [Data Sources and Exploratory Data Analysis](#)

establishing business objectives, [Establishing Business Objectives](#)

evaluating and comparing models, [Evaluating and Comparing Models-Impact of Responsible AI on Modeling](#)

experimentation, [Experimentation-Experimentation](#)

feature engineering and selection, [Feature Engineering and Selection](#), [Feature Engineering and Selection-How Feature Selection Impacts MLOps Strategy](#)

in practice, [In Practice](#)

in theory, [What Is a Machine Learning Model?](#)

MLOps considerations by algorithm type, [Different ML Algorithms](#), [Different MLOps Challenges](#)

required components, [Required Components](#)

Responsible AI, [Responsible AI](#)

training and evaluation, [Training and Evaluation](#)

version management and reproducibility, [Version Management and Reproducibility-Version Management and Reproducibility](#)

iterations and life cycle, [Iteration and Life Cycle-The Feedback Loop](#)

life cycle, [Defining MLOps and Its Challenges](#)

mitigating risks with MLOps, [MLOps to Mitigate Risk-MLOps for Responsible AI](#)

monitoring after deployment to production, [Monitoring-Business Concerns](#)

productionalization and deployment, [Productionalization and Deployment-Model Deployment Requirements](#)

monitoring and refining governance, [Step 8: Monitor and Refine](#)

monitoring machine learning models, [Monitoring-Business Concerns](#), [Monitoring and Feedback Loop-Closing Thoughts](#)

business concerns, [Business Concerns](#)

consumer credit risk management model, [Prepare for Production, Deploy to Production](#)

consumption forecast for power grid model, [Monitoring](#)

data scientist concerns, [Data Scientists](#), [Data Scientist Concerns](#)

deciding how often to retrain models, [How Often Should Models Be Retrained?-How Often Should Models Be Retrained?](#)

degradation of performance, understanding, [Understanding Model Degradation-Input Drift Detection](#)

ground truth evaluation, [Ground Truth Evaluation](#)

input drift detection, [Input Drift Detection](#)

DevOps concerns, [DevOps Concerns](#)

drift detection in practice, [Drift Detection in Practice-Interpretation of results](#)

input drift detection techniques, [Input Drift Detection Techniques](#)

feedback loop, [The Feedback Loop-Closing Thoughts](#)

logging, [Logging](#)

model evaluation, [Model Evaluation-Model evaluation store](#)

online evaluation of models in production, [Online Evaluation-A/B testing](#)

marketing recommendation engine, [Monitoring and Retraining Strategy](#), [Monitoring and Feedback-Monitoring Performance](#)

model runs overnight, sleeps in daytime, [Runs Overnight, Sleeps During Daytime](#)  
monitoring performance, [Monitoring Performance](#)  
option to automatically control models, [Option to Automatically Control Models](#)  
option to manually control models, [Option to Manually Control Models](#)  
retraining models, [Retraining Models](#)  
updating models, [Updating Models](#)  
model risk manager/auditor, role in, [Model Risk Manager/Auditor in production](#), [Maintenance in Production](#)  
monotonicity constraints, [Model Development](#)  
MRM (model risk management), [Model Risk Manager/Auditor](#) (see also model risk manager/auditor)  
principles for good MRM, defined by UK PRA, [Financial Model Risk Management Regulation](#)  
regulation for financial models, [Financial Model Risk Management Regulation](#)  
multi-armed bandit testing, [The Feedback Loop](#)

## N

neural networks, [Different ML Algorithms](#), [Different MLOps Challenges](#)  
adversarial attacks on, [Adversarial Attacks](#)  
long short-term memory (LSTM), [Modeling](#)  
NLP (natural language processing) pretrained models, compression used with, [Performance considerations](#)  
null hypothesis to p-values, [Ground Truth Evaluation](#)

## O

one-hot encoding, [Feature Engineering Techniques](#)  
online evaluation of models in production, [Online Evaluation-A/B testing](#)  
A/B testing, [A/B testing](#)  
champion/challenger, [Champion/Challenger](#)

online machine learning, [How Often Should Models Be Retrained?](#)  
operationalization and MLOps, [Data Scientists](#)  
orchestration of containers, [Model Deployment Types and Contents](#)

## P

p-values, [Ground Truth Evaluation](#)

advantages and drawbacks of, [Univariate statistical tests](#)

parallelization of batch scoring, [Categories of Model Deployment](#)

partitioning (or sharding), [Categories of Model Deployment](#)

distributing batch processing by partitioning data, [Scaling](#)

[Deployments](#)

people in MLOps, roles and responsibilities, [People of MLOps-Closing Thoughts](#)

data engineers, [Data Engineers](#)

data scientists, [Data Scientists-Data Scientists](#)

DevOps, [DevOps](#)

machine learning architects, [Machine Learning Architect](#)

model risk manager/auditor, [Model Risk Manager/Auditor](#)

overview, [People of MLOps](#)

software engineers, [Software Engineers](#)

subject matter experts (SMEs), [Subject Matter Experts-Subject Matter Experts](#)

performance

considerations when converting from development to production environments, [Performance considerations](#)

degradation of model performance, monitoring, [Understanding Model Degradation-Input Drift Detection](#)

drift, assessing for subpopulations in credit risk management model, [Prepare for Production](#)

model retraining and, [How Often Should Models Be Retrained?](#)

monitoring for marketing recommendation engines, [Monitoring Performance](#)

monitoring ML models for, [Monitoring and Feedback Loop](#)

testing on out-of-sample datasets for consumer credit risk management model, [Model Bias Considerations](#)

performance metrics, [Required Components](#)

personally identifiable information (PII), [Data Sources and Exploratory Data Analysis](#)

making sure no PII used to train a model, [Data Governance](#)  
pharmaceutical regulations in US, GxP, [Pharmaceutical Regulation in the US: GxP](#)

poisoning attacks, [Adversarial Attacks](#)

portable formats for models, [Model Deployment Types and Contents, Deployment](#)

power systems, [Power Systems-Power Systems](#)

predictions

conformal, [Model Misbehavior](#)

decoupled ground truth and prediction, [Ground Truth Evaluation](#)

examples of model prediction and generalization, [In Theory](#)

ground truth for model predictions, [Ground truth](#)

input drift and, [Input drift](#)

probabilities

distances between probability distributions, measuring, [Prepare for Production](#)

produced by models used for risk assessment, [Model Bias Considerations](#)

problem definition and data acquisition, consumer credit risk management MLOps application, [Background: The Business Use Case](#)

problem definition, using machine learning or not, [Problem Definition: Machine Learning, or Not Machine Learning?-Spatial and Temporal Resolution](#)

process governance, [Process Governance](#)

effective implementation, difficulties of, [Process Governance](#)

production, deploying to, [Deploying to Production-Closing Thoughts](#)

building ML artifacts, [Building ML Artifacts-The Testing Pipeline](#)

using testing pipeline on model, [The Testing Pipeline](#)

CI/CD pipelines, [CI/CD Pipelines](#)

consumer credit risk management model, [Deploy to Production](#)

consumption forecast prototype models, [Deployment](#)

containerization, [Containerization-Containerization](#)

deployment strategies, [Deployment Strategies-Maintenance in Production](#)

categories of model deployment, [Categories of Model Deployment](#)

considerations in sending models to production, [Considerations When Sending Models to Production](#)

maintenance of models in production, [Maintenance in Production](#)

scaling deployments, [Scaling Deployments-Requirements and Challenges](#)

requirements and challenges, [Requirements and Challenges](#)

production, preparing for, [Preparing for Production-Closing Thoughts](#)

consumer credit risk management model, [Prepare for Production](#)

key ideas, summary of, [Closing Thoughts](#)

machine learning security, [Machine Learning Security-Other Vulnerabilities](#)

model risk evaluation, [Model Risk Evaluation-The Origins of ML Model Risk](#)

model risk mitigation, [Model Risk Mitigation-Model Misbehavior](#)

changing environments, [Changing Environments](#)

interactions between models, [Interactions Between Models](#)

model misbehavior, [Model Misbehavior](#)

quality assurance for machine learning, [Quality Assurance for Machine Learning-Key Testing Considerations](#)

reproducibility and auditability for models, [Reproducibility and Auditability](#)

runtime environments, [Runtime Environments-Final Thoughts on Runtime Environments](#)

adaptation from development to production environments, [Adaptation from Development to Production Environments-Performance considerations](#)

data access before validation and launch to production, [Data Access Before Validation and Launch to Production](#)

final thoughts on, [Final Thoughts on Runtime Environments](#)

productionalization and deployment of models, [Productionalization and Deployment-Model Deployment Requirements](#)

model deployment requirements, [Model Deployment Requirements](#)

model deployment types and contents, [Model Deployment Types and Contents](#)



progressive or canary rollouts, [Model Risk Mitigation](#)  
project criticality and operationalization approaches to risk assessment, [Matching Governance with Risk Level](#)  
provenance of data, [Element 1: Data](#)  
pruning models, [Performance considerations](#)  
public opinion, influence on ML governance, [Who Decides What Governance the Organization Needs?](#)  
push or pull recommendations, [Push or Pull?](#)  
Python, [Risk Mitigation](#), [Tooling considerations](#), [Deployment](#)

## Q

QA (quality assurance) for machine learning, [Quality Assurance for Machine Learning-Key Testing Considerations](#)  
key testing considerations, [Key Testing Considerations](#)  
providing clear view of model performance and facilitating auditability, [Reproducibility and Auditability](#)  
quantization, [Performance considerations](#)

## R

R language, [Deployment](#)  
RACI (responsible, accountable, consulted, informed), [Step 3: Establish Responsibilities](#)  
randomness, [Version Management and Reproducibility](#)  
random sampling, [Ground Truth Evaluation](#)  
real-time scoring, [Categories of Model Deployment](#)  
logging streaming data, [Requirements and Challenges](#)  
marketing recommendation engine model, [Real-Time Scoring](#)  
recommendation engines, [MLOps in Practice: Marketing Recommendation Engines](#)  
(see also marketing recommendation engines)  
rise of, [The Rise of Recommendation Engines-Push or Pull?](#)  
deciding on push or pull recommendations, [Push or Pull?](#)  
turning recommendations on/off, [Ability to Turn Recommendations On and Off](#)

red-black deployment, [Considerations When Sending Models to Production](#)

regulations, [Step 1: Understand and Classify the Analytics Use Cases](#)

AI-specific, new wave of, [The New Wave of AI-Specific Regulations-The New Wave of AI-Specific Regulations](#)

current, driving MLOps governance, [Current Regulations Driving MLOps Governance-GDPR and CCPA Data Privacy Regulations](#)

financial model risk management, [Financial Model Risk Management Regulation](#)

GDPR and CCPA data privacy regulations, [GDPR and CCPA Data Privacy Regulations](#)

pharmaceutical regulation in US, GxP, [Pharmaceutical Regulation in the US: GxP](#)

governance and regulatory checks in model deployments, [Model Deployment Requirements](#)

government regulations on ML to mitigate negative impact of its use, [Governance](#)

government regulations on use of personal data by businesses, [Governance](#)

process governance and, [Process Governance](#)

releases, [Deployment Strategies](#)

reproducibility

auditability and, [Reproducibility and Auditability](#)

of experiments with ML models, [Reproducibility](#)

importance as model property, [Version Management and Reproducibility](#)

in MLOps versus academia, [Reproducibility and Auditability](#)

resource demands of models

capped, [Model Deployment Requirements](#)

computing power, [Different ML Algorithms, Different MLOps Challenges](#)

scalability of compute resources for deployed models, [DevOps Concerns](#)

resource monitoring, [Maintenance in Production](#)

resources

monitoring for ML models, [Monitoring and Feedback Loop](#)

optimizing usage with elastic systems, [Scaling Deployments](#)

responsibilities, establishing in MLOps, [Step 3: Establish Responsibilities](#)

Responsible AI, [The Emergence of Responsible AI-Element 5: Governance](#)

businesses engaging with, [Governance](#)

impacts on modeling, [Impact of Responsible AI on Modeling](#)

key elements of, [Key Elements of Responsible AI-Element 5: Governance](#)

bias, [Element 2: Bias](#)

data, [Element 1: Data](#)

governance, [Element 5: Governance](#)

inclusiveness, [Element 3: Inclusiveness](#)

model management at scale, [Element 4: Model Management at Scale](#)

in marketing recommendation engine, [Ability to Turn Recommendations On and Off](#)

in ML model development, [Responsible AI](#)

MLOps for, [MLOps for Responsible AI](#)

model explainability, [Responsible AI](#)

results of in-depth analysis of models, comparing, [Version Management and Reproducibility](#)

retraining models

deciding how often to retrain, [How Often Should Models Be Retrained?-How Often Should Models Be Retrained?](#)

marketing recommendation engine, [Monitoring and Retraining Strategy, Retraining Models](#)

reweighting, [Drift Detection in Practice](#)

for biased samples, [Ground Truth Evaluation, Example Causes of Data Drift](#)

(see also biases)

in consumer credit risk management model bias, [Model Bias Considerations](#)

feature that causes drift, [Interpretation of results](#)

risks

assessing, considerations in MLOps, [Matching Governance with Risk Level](#)

credit risk modeling, [Model Development](#)

(see also consumer credit risk management)

financial model risk management regulation, [Financial Model Risk Management Regulation](#)

matching governance with risk level, [Matching Governance with Risk Level](#)

model risk evaluation, [Model Risk Evaluation-The Origins of ML Model Risk](#)

origins of ML model risk, [The Origins of ML Model Risk](#)

purpose of model validation, [The Purpose of Model Validation](#)

model risk manager/auditor, [Model Risk Manager/Auditor](#)

model risk mitigation, [Model Risk Mitigation-Model Misbehavior](#)

changing environments, [Changing Environments](#)

interactions between models, [Interactions Between Models](#)

model misbehavior, [Model Misbehavior](#)

need for proactively addressing in ML, [Who Decides What Governance the Organization Needs?](#)

risk assessment, [Risk Assessment](#)

risk mitigation, MLOps for, [Risk Mitigation](#)

RMSE (root-mean-square error), [Updating Models](#)

rollbacks to previous model versions, [Maintenance in Production](#)

runtime environments, [Runtime Environments-Final Thoughts on Runtime Environments](#)

adaptation from development to production environments, [Adaptation from Development to Production Environments-Performance considerations](#)

performance considerations, [Performance considerations](#)

tooling considerations, [Tooling considerations](#)

data access before validation and launch to production, [Data Access Before Validation and Launch to Production](#)

final thoughts on, [Final Thoughts on Runtime Environments](#)

## S

sampling

random, [Ground Truth Evaluation](#)

selection bias, [Example Causes of Data Drift](#)

SCADA (supervisory control and data acquisition) system, [Data Collection](#)

scale

MLOps for, [MLOps for Scale](#)

model management at scale (Responsible AI), [Element 4: Model Management at Scale](#)

scalability of compute resources for deployed models, [DevOps Concerns](#)

scalability of marketing recommendation engine model, [Scalability and Customizability](#)

scaling deployments, [Scaling Deployments-Requirements and Challenges](#)

requirements and challenges, [Requirements and Challenges](#)

scalable and elastic systems, [Scaling Deployments](#)

scaling the number of models, [Scaling Deployments](#)

scaling inference on models, [Final Thoughts on Runtime Environments](#)

scikit-learn, [Risk Mitigation](#)

models developed with, adaptation to production environments, [Tooling considerations](#)

security in machine learning, [Machine Learning Security-Other Vulnerabilities](#)

adversarial attacks, [Adversarial Attacks](#)

other vulnerabilities, [Other Vulnerabilities](#)

selection bias, [Example Causes of Data Drift](#)

settings, reproducibility and, [Version Management and Reproducibility](#)

shadow testing, [Online Evaluation](#)

(see also champion/challenger)

of new model version deployed alongside existing model, [The Feedback Loop](#)

Shapley values, [Responsible AI](#), [Impact of Responsible AI on Modeling, Model Development](#)

sharding, [Categories of Model Deployment](#)

SMEs (see subject matter experts)

software engineers, [Software Engineers](#)

role in and needs from MLOps, [Software Engineers](#)

role in machine learning model life cycle, [Software Engineers](#)  
Spark, [Scaling Deployments](#)  
spatial aggregations of data, [Implementation](#)  
spatial and temporal resolution, forecast uncertainty and, [Spatial and Temporal Resolution](#)  
statistics  
from null hypothesis to p-values, [Ground Truth Evaluation](#)  
statistical metrics for model performance monitoring, [Ground Truth Evaluation](#)  
statistical results on model testing, [Key Testing Considerations](#)  
statistically driven decision-making processes in ML, [Who Decides What Governance the Organization Needs?](#)  
univariate statistical tests, [Input Drift Detection Techniques](#)  
stochasticity, [Modeling](#)  
subject matter experts (SMEs), [Subject Matter Experts-Subject Matter Experts](#)  
role in and needs from MLOps, [Subject Matter Experts](#)  
role in machine learning model life cycle, [Subject Matter Experts](#)  
subpopulations  
analyses of, [Responsible AI](#)  
analysis and model fairness, [Key Testing Considerations](#)  
investing model performance for in consumer credit risk management, [Model Bias Considerations](#)  
statistical tests on results in preproduction model testing, [Key Testing Considerations](#)  
supervised learning, [Required Components](#)

## T

technical debt in machine learning, [How Feature Selection Impacts MLOps Strategy](#)  
temporal aggregations of data, [Implementation](#)  
testing  
key testing considerations, [Key Testing Considerations](#)  
for models using self-training algorithm, [How Often Should Models Be Retrained?](#)  
testing pipeline for ML artifacts, [The Testing Pipeline](#)

tools

considerations in adaptation from development to production environments, [Tooling considerations](#)

selecting for centralized governance management, [Step 6: Select the Tools for Centralized Governance Management](#)

traceability in pharmaceutical industry, [Pharmaceutical Regulation in the US: GxP](#)

training data, [A Primer on Machine Learning](#), [Required Components](#)

adversarial attacks on, [Adversarial Attacks](#)

data governance concerns, [Data Governance](#)

keeping up to date for deployed models, [Data Scientist Concerns](#)

non-stationary environment, [Example Causes of Data Drift](#)

quality of, determining model performance, [Monitoring and Feedback Loop](#)

training models, [Training and Evaluation](#)

automation in, [Experimentation](#)

deciding how often to retrain models, [How Often Should Models Be Retrained?-How Often Should Models Be Retrained?](#)

marketing recommendation engine model, [Model Training and Deployment-Ability to Turn Recommendations On and Off](#)

retraining existing model with latest training data, [Iteration](#)

transfer learning, [Feature Engineering Techniques](#)

transparency

data scientists' need for from MLOps, [Data Scientists](#)

subject matter experts, role in MLOps, [Subject Matter Experts](#)

transparent strategies for machine learning, [Closing Thoughts](#)

tree-based algorithms, [Different ML Algorithms](#), [Different MLOps Challenges](#)

trust

EU's requirements for trustworthy AI applications, [The New Wave of AI-Specific Regulations](#)

importance to consumers and businesses, [Governance](#)

Tweedie distribution, [Model Development](#)

Twitter chatbot by Microsoft, [Adversarial Attacks](#)

two-stage models with offset, [Model Development](#)



U.S. Food and Drug Administration (FDA), [Pharmaceutical Regulation in the US: GxP](#)

UCI Wine Quality dataset, [Input Drift Detection](#)

UK Prudential Regulation Authority's (PRA) regulation, [Financial Model Risk Management Regulation](#)

underfitting, [Experimentation](#)

univariate statistical tests, [Input Drift Detection Techniques](#)

updating models, [Updating Models](#)

## V

validation

data access before model validation, [Data Access Before Validation and Launch to Production](#)

model for consumer credit risk management, [Model Bias Considerations](#)

purpose of model validation, [The Purpose of Model Validation](#)

testing on recent production data from models, [The Testing Pipeline](#)

variance in models, [Experimentation](#)

version control systems, central, [Building ML Artifacts](#)

version management, [Containerization](#)

version management and reproducibility, [Reproducibility, Version Management and Reproducibility-Version Management and Reproducibility](#)

VMs (virtual machines), containers versus, [Containerization](#)

## W

Wasserstein distance, [Prepare for Production](#)

what-if analysis, [Responsible AI](#)

## X

XGBoost algorithm, [Model Development](#)

[Support](#)   [Sign Out](#)

©2022 O'REILLY MEDIA, INC. [TERMS OF SERVICE](#) [PRIVACY POLICY](#)